



Logistic Regression (with Pytorch)

김재광 교수 (소프트웨어융합대학 글로벌융합학부)



Contents

- Reminder
- Imports
- Training Data
- Computing Hypothesis
- Computing Cost Function
- Evaluation
- Higher Implementation



Reminder

- Hypothesis

$$H(x) = \frac{1}{1 + e^{-W^T X}}$$

- Cost

$$\text{cost}(W) = -\frac{1}{m} \sum_{i=1}^m y \log(H(x)) + (1 - y)(\log(1 - H(x)))$$

- If $y \cong H(x)$, cost is near 0
- If $y \neq H(x)$, cost is high

Reminder (Cont'd)

- Weight update via Gradient Descent

$$\begin{aligned} W &:= W - \alpha \nabla W \\ &= W - \alpha \frac{\partial}{\partial W} \text{cost}(W) \end{aligned}$$

Imports

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
```

```
# for reproducibility
torch.manual_seed(1)
```

Training Data

```
x_data = [[1, 2], [2, 3], [3, 1], [4, 3], [5, 3], [6, 2]]  
y_data = [[0], [0], [0], [1], [1], [1]]
```

```
x_train = torch.FloatTensor(x_data)  
y_train = torch.FloatTensor(y_data)
```

```
print(x_train.shape)  
print(y_train.shape)
```

```
torch.Size([6, 2])  
torch.Size([6, 1])
```



Computing Hypothesis

$$H(x) = \frac{1}{1 + e^{-W^T X}}$$

```
print('e^1 equals: ', torch.exp(torch.FloatTensor([1])))
```

```
W = torch.zeros((2, 1), requires_grad=True)  
b = torch.zeros(1, requires_grad=True)
```

```
hypothesis = 1 / (1 + torch.exp(-(x_train.matmul(W) + b)))
```

```
print(hypothesis)  
print(hypothesis.shape)
```

```
tensor([[0.5000],  
        [0.5000],  
        [0.5000],  
        [0.5000],  
        [0.5000],  
        [0.5000]],  
       grad_fn=<MulBackward>)  
torch.Size([6, 1])
```



Computing Hypothesis (Cont'd)

```
print('1/(1+e^{-1}) equals: ', torch.sigmoid(torch.FloatTensor([1])))
```

```
hypothesis = torch.sigmoid(x_train.matmul(W) + b)
```

```
print(hypothesis)  
print(hypothesis.shape)
```

```
tensor([[0.5000],  
        [0.5000],  
        [0.5000],  
        [0.5000],  
        [0.5000],  
        [0.5000]],  
       grad_fn=<SigmoidBackward>)  
torch.Size([6, 1])
```




Computing Cost Function

$$\text{cost}(W) = -\frac{1}{m} \sum_{i=1}^m y \log(H(x)) + (1 - y) (\log(1 - H(x)))$$

```
print(hypothesis)
print(y_train)
```

```
tensor([[0.5000],
        [0.5000],
        [0.5000],
        [0.5000],
        [0.5000],
        [0.5000]], grad_fn=<SigmoidBackward>)
tensor([[0.],
        [0.],
        [0.],
        [1.],
        [1.],
        [1.]])
```


Computing Cost Function (Cont'd)

$$\text{cost}(W) = -\frac{1}{m} \sum_{i=1}^m y \log(H(x)) + (1 - y) (\log(1 - H(x)))$$

```
cost = losses.mean()  
print(cost)
```

```
tensor(0.6931, grad_fn=<MeanBackward1>)
```

Computing Cost Function (Cont'd)

- with `F.binary_cross_entropy`

```
F.binary_cross_entropy(hypothesis, y_train)
```

```
tensor(0.6931, grad_fn=<BinaryCrossEntropyBackward>)
```

Computing Cost Function (Cont'd)

모델 초기화

```
W = torch.zeros((2, 1), requires_grad=True)
```

```
b = torch.zeros(1, requires_grad=True)
```

optimizer 설정

```
optimizer = optim.SGD([W, b], lr=1)
```

Computing Cost Function (Cont'd)

```
nb_epochs = 1000
for epoch in range(nb_epochs + 1):

    # Cost 계산
    hypothesis = torch.sigmoid(x_train.matmul(W) + b) # or .mm or @
    cost = -(y_train * torch.log(hypothesis) +
             (1 - y_train) * torch.log(1 - hypothesis)).mean()

    # cost로 H(x) 개선
    optimizer.zero_grad()
    cost.backward()
    optimizer.step()

    # 100번마다 로그 출력
    if epoch % 100 == 0:
        print('Epoch {:4d}/{:4d} Cost: {:.6f}'.format(
            epoch, nb_epochs, cost.item()
        ))
```

Epoch	0/1000	Cost :	0.693147
Epoch	100/1000	Cost :	0.134722
Epoch	200/1000	Cost :	0.080643
Epoch	300/1000	Cost :	0.057900
Epoch	400/1000	Cost :	0.045300
Epoch	500/1000	Cost :	0.037261
Epoch	600/1000	Cost :	0.031673
Epoch	700/1000	Cost :	0.027556
Epoch	800/1000	Cost :	0.024394
Epoch	900/1000	Cost :	0.021888
Epoch	1000/1000	Cost :	0.019852

Computing Cost Function (Cont'd)

```
nb_epochs = 1000
for epoch in range(nb_epochs + 1):

    # Cost 계산
    hypothesis = torch.sigmoid(x_train.matmul(W) + b) # or .mm or @
    cost = F.binary_cross_entropy(hypothesis, y_train)

    # cost로 H(x) 개선
    optimizer.zero_grad()
    cost.backward()
    optimizer.step()

    # 100번마다 로그 출력
    if epoch % 100 == 0:
        print('Epoch {:4d}/{:4d} Cost: {:.6f}'.format(
            epoch, nb_epochs, cost.item()
        ))
```

Epoch	0/1000	Cost :	0.693147
Epoch	100/1000	Cost :	0.134722
Epoch	200/1000	Cost :	0.080643
Epoch	300/1000	Cost :	0.057900
Epoch	400/1000	Cost :	0.045300
Epoch	500/1000	Cost :	0.037261
Epoch	600/1000	Cost :	0.031673
Epoch	700/1000	Cost :	0.027556
Epoch	800/1000	Cost :	0.024394
Epoch	900/1000	Cost :	0.021888
Epoch	1000/1000	Cost :	0.019852

Evaluation

```
hypothesis = torch.sigmoid(x_train.matmul(W) + b)
print(hypothesis[:5])
```

```
tensor([[0.4103],
        [0.9242],
        [0.2300],
        [0.9411],
        [0.1772]], grad_fn=<SliceBackward>)
```

```
prediction = hypothesis >= torch.FloatTensor([0.5])
print(prediction[:5])
```

```
tensor([[0],
        [1],
        [0],
        [1],
        [0]], dtype=torch.uint8)
```




Evaluation (Cont'd)

```
print(prediction[:5])  
print(y_train[:5])
```

```
tensor([[0],  
        [1],  
        [0],  
        [1],  
        [0]], dtype=torch.uint8)  
tensor([[0.],  
        [1.],  
        [0.],  
        [1.],  
        [0.]])
```



Evaluation (Cont'd)

```
correct_prediction = prediction.float() == y_train
print(correct_prediction[:5])
```

```
tensor([[1],
        [1],
        [1],
        [1],
        [1]], dtype=torch.uint8)
```

```
accuracy = correct_prediction.sum().item() / len(correct_prediction)
print('The accuracy of {:.2f}% for the training set.'.format(accuracy * 100))
```

The accuracy of 76.68% for the training set.



Higher Implementation

```
class BinaryClassifier(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.linear = nn.Linear(8, 1)  
        self.sigmoid = nn.Sigmoid()  
  
    def forward(self, x):  
        return self.sigmoid(self.linear(x))
```

```
model = BinaryClassifier()  
  
# optimizer 설정  
optimizer = optim.SGD(model.parameters(), lr=1)  
  
nb_epochs = 100
```

Higher Implementation (Cont'd)

```
for epoch in range(nb_epochs + 1):
```

```
    # H(x) 계산
```

```
    hypothesis = model(x_train)
```

```
    # cost 계산
```

```
    cost = F.binary_cross_entropy(hypothesis, y_train)
```

```
    # cost로 H(x) 개선
```

```
    optimizer.zero_grad()
```

```
    cost.backward()
```

```
    optimizer.step()
```

```
    # 10번마다 로그 출력
```

```
    if epoch % 10 == 0:
```

```
        prediction = hypothesis >= torch.FloatTensor([0.5])
```

```
        correct_prediction = prediction.float() == y_train
```

```
        accuracy = correct_prediction.sum().item() / len(correct_prediction)
```

```
        print('Epoch {:4d}/{:4d} Cost: {:.6f} Accuracy {:.2f}%'.format(
            epoch, nb_epochs, cost.item(), accuracy * 100,
        ))
```

Epoch	0/100	Cost :	0.704829	Accuracy	45.72%
Epoch	10/100	Cost :	0.572391	Accuracy	67.59%
Epoch	20/100	Cost :	0.539563	Accuracy	73.25%
Epoch	30/100	Cost :	0.520042	Accuracy	75.89%
Epoch	40/100	Cost :	0.507561	Accuracy	76.15%
Epoch	50/100	Cost :	0.499125	Accuracy	76.42%
Epoch	60/100	Cost :	0.493177	Accuracy	77.21%
Epoch	70/100	Cost :	0.488846	Accuracy	76.81%
Epoch	80/100	Cost :	0.485612	Accuracy	76.28%
Epoch	90/100	Cost :	0.483146	Accuracy	76.55%
Epoch	100/100	Cost :	0.481234	Accuracy	76.81%