


MNIST with CNN (with Pytorch)


김재광 교수 (소프트웨어융합대학 글로벌융합학부)



C

ontents

- What is MNIST?
- MNIST Data Manipulation
- CNN 학습 단계
- MNIST CNN code



What is MNIST?

- MNIST: handwritten digits dataset



- train-images-idx3-ubyte.gz: training set images (9912422 bytes)
- train-labels-idx1-ubyte.gz: training set labels (28881 bytes)
- t10k-images-idx3-ubyte.gz: test set images (1648877 bytes)
- t10k-labels-idx1-ubyte.gz: test set labels (4542 bytes)



What is MNIST?(Cont'd)

● Example of MNIST


- 28x28 image
- 1 channel gray image
- 0-9 digits

for X, Y in data_loader:

reshape input image into [batch_size by 784]

Label is not one-hot encoded

X = X.view(-1, 28*28)





What is MNIST?(Cont'd)

● TORCHVISION

- The TORCHVISION package consists of popular datasets, model architectures, and common image transformations for computer vision.

- torchvision.datasets
 - MNIST
 - Fashion-MNIST
 - KMNIST
 - EMNIST
 - QMNIST
 - FakeData
 - COCO
 - LSUN
 - ImageFolder
 - DatasetFolder
 - ImageNet
- CIFAR
- STL10
- SVHN
- PhotoTour
- SBU
- Flickr
- VOC
- Cityscapes
- SBD
- USPS
- Kinetics-400
- HMDB51
- UCF101
- CelebA
- torchvision.io
 - Video
- torchvision.models
 - Classification
 - Semantic Segmentation
 - Object Detection, Instance Segmentation and Person Keypoint Detection
 - Video classification
- torchvision.ops
- torchvision.transforms
 - Transforms on PIL Image
 - Transforms on torch.*Tensor
 - Conversion Transforms
 - Generic Transforms
 - Functional Transforms
- torchvision.utils



<https://pytorch.org/docs/stable/torchvision/index.html>



MNIST Data Manipulation

● Reading data

```
import torchvision.datasets as dsets
...
mnist_train = dsets.MNIST(root="MNIST_data/", train=True, transform=transforms.ToTensor(),download=True)
mnist_test = dsets.MNIST(root="MNIST_data/", train=False, transform=transforms.ToTensor(),download=True)
data_loader = torch.utils.DataLoader(DataLoader=mnist_train, batch_size=batch_size,shuffle=True, drop_last=True)
...
for epoch in range(training_epochs):
    ...
    for X, Y in data_loader:
        # reshape input image into [batch_size by 784]
        # label is not one-hot encoded
        X = X.view(-1, 28 * 28).to(device)
```

MNIST Data Manipulation (Cont'd)

- **Epoch / Batch size / Iteration**

- **One epoch**

- ✓ One forward / backward pass of all the training examples

- **Batch size**

- ✓ # of training examples in one forward / backward pass
- ✓ The higher the batch size, the more memory space is required

- **# of Iteration**

- ✓ # of passes, each pass using [batch size] # of examples To be clear, one pass = one forward pass + one backward pass (we do not count the forward pass and backward pass as two different passes)

ex

- ❖ if you have 1000 training examples, and your batch size is 500, then it will take 2 iterations to complete 1 epoch

MNIST Data Manipulation (Cont'd)

● Softmax

```
# MNIST data image of shape 28 * 28 = 784 Softmax
linear = torch.nn.Linear(784, 10, bias=True).to(device)
# initialization
torch.nn.init.normal_(linear.weight)
# parameters
training_epochs = 15
batch_size = 100
# define cost/loss & optimizer
criterion = torch.nn.CrossEntropyLoss().to(device) # Softmax is internally computed.
optimizer = torch.optim.SGD(linear.parameters(), lr=0.1)
```


MNIST Data Manipulation (Cont'd)

● Softmax (cont'd)

```
for epoch in range(training_epochs):
    avg_cost = 0
    total_batch = len(data_loader)
    for X, Y in data_loader:
        # reshape input image into [batch_size by 784]
        # label is not one-hot encoded
        X = X.view(-1, 28 * 28).to(device)
        optimier.zero_grad()
        hypothesis = linear(X)
        cost = criterion(hypothesis, Y)
        cost.backward()
        avg_cost += cost / total_batch

    print("Epoch: ", "%04d" % (epoch+1), "cost =", "{:.9f}".format(avg_cost))
```

MNIST Data Manipulation (Cont'd)

- Softmax (cont'd)

```
Epoch: 0001 cost = 2.511683702
Epoch: 0002 cost = 0.977319956
Epoch: 0003 cost = 0.797017217
Epoch: 0004 cost = 0.710427940
Epoch: 0005 cost = 0.655205429
Epoch: 0006 cost = 0.615207732
Epoch: 0007 cost = 0.584421575
Epoch: 0008 cost = 0.559486568
Epoch: 0009 cost = 0.538655698
Epoch: 0010 cost = 0.520880997
Epoch: 0011 cost = 0.505315244
Epoch: 0012 cost = 0.491431117
Epoch: 0013 cost = 0.479477882
Epoch: 0014 cost = 0.468681127
Epoch: 0015 cost = 0.458788306
Learning finished
Accuracy: 0.8718999624252319
```

MNIST Data Manipulation (Cont'd)

● Test

```
# Test the model using test sets
```

```
With torch.no_grad():
```

```
    X_test = mnist_test.test_data.view(-1, 28 * 28).float().to(device)
```

```
    Y_test = mnist_test.test_labels.to(device)
```

```
    prediction = linear(X_test)
```

```
    correct_prediction = torch.argmax(prediction, 1) == Y_test
```

```
    accuracy = correct_prediction.float().mean()
```

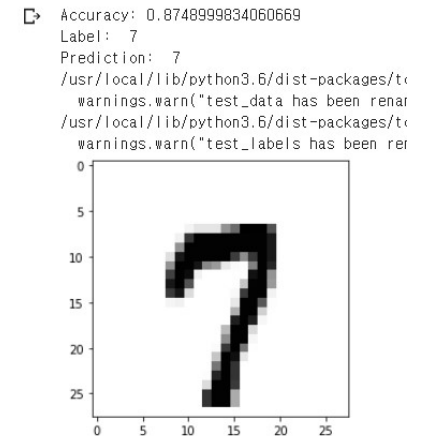
```
    print("Accuracy: ", accuracy.item())
```

MNIST Data Manipulation (Cont'd)

● Visualization

```
import matplotlib.pyplot as plt
import random

...
r = random.randint(0, len(mnist_test) - 1)
X_single_data = mnist_test.test_data[r:r + 1].view(-1, 28 * 28).float().to(device)
Y_single_data = mnist_test.test_labels[r:r + 1].to(device)
print("Label: ", Y_single_data.item())
single_prediction = linear(X_single_data)
print("Prediction: ", torch.argmax(single_prediction, 1).item())
plt.imshow(mnist_test.test_data[r:r + 1].view(28, 28), cmap="Greys", interpolation="nearest")
plt.show()
```



CNN 학습 단계

- 라이브러리를 가져오기
- GPU 사용 설정 및 random seed 설정
- 학습에 필요한 parameter 설정
 - lr, training epochs, batch_size, ...
- 데이터 셋 로드
- 학습 모델 생성
- Loss function 및 optimizer 설정
- Loss 확인을 통해 모델 학습
- 성능 확인



MNIST CNN Code

```
import torch
import torchvision.datasets as dsets
import torchvision.transforms as transforms
import torch.nn.init

device = 'cuda' if torch.cuda.is_available() else 'cpu'


# for reproducibility
torch.manual_seed(777)
if device == 'cuda':
    torch.cuda.manual_seed_all(777)
```




MNIST CNN Code (Cont'd)

```
# CNN Model (2 conv layers)
class CNN(torch.nn.Module):

    def __init__(self):
        super(CNN, self).__init__()
        # L1 ImgIn shape=(?, 28, 28, 1)
        #  Conv  -> (?, 28, 28, 32)
        #  Pool  -> (?, 14, 14, 32)
        self.layer1 = torch.nn.Sequential(
            torch.nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2))
        # L2 ImgIn shape=(?, 14, 14, 32)
        #  Conv  -> (?, 14, 14, 64)
        #  Pool  -> (?, 7, 7, 64)
        self.layer2 = torch.nn.Sequential(
            torch.nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2))
        # Final FC 7x7x64 inputs -> 10 outputs
        self.fc = torch.nn.Linear(7 * 7 * 64, 10, bias=True)
        torch.nn.init.xavier_uniform_(self.fc.weight)
```





MNIST CNN Code (Cont'd)

```
def forward(self, x):  
    out = self.layer1(x)  
    out = self.layer2(out)  
    out = out.view(out.size(0), -1) # Flatten them for FC  
    out = self.fc(out)  
    return out
```

```
# instantiate CNN model  
model = CNN().to(device)  
  
# define cost/loss & optimizer  
criterion = torch.nn.CrossEntropyLoss().to(device) # Softmax is internally computed.  
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```



MNIST CNN Code (Cont'd)

```
# train my model
total_batch = len(data_loader)
print('Learning started. It takes sometime.')
for epoch in range(training_epochs):
    avg_cost = 0
    for X, Y in data_loader:
        # image is already size of (28x28), no reshape
        # label is not one-hot encoded
        X = X.to(device)
        Y = Y.to(device)

        optimizer.zero_grad()
        hypothesis = model(X)
        cost = criterion(hypothesis, Y)
        cost.backward()
        optimizer.step()

    avg_cost += cost / total_batch
    print('[Epoch: {:>4}] cost = {:>.9}'.format(epoch + 1, avg_cost))
print('Learning Finished!')
```

➡ Learning started. It takes sometime.

[Epoch: 1]	cost = 0.159776822
[Epoch: 2]	cost = 0.0413324498
[Epoch: 3]	cost = 0.0292899609
[Epoch: 4]	cost = 0.02193336
[Epoch: 5]	cost = 0.0172540527
[Epoch: 6]	cost = 0.0135775162
[Epoch: 7]	cost = 0.0111167822
[Epoch: 8]	cost = 0.0123288101
[Epoch: 9]	cost = 0.0087507898
[Epoch: 10]	cost = 0.0086597465
[Epoch: 11]	cost = 0.00723933568
[Epoch: 12]	cost = 0.00641156221
[Epoch: 13]	cost = 0.00667931139
[Epoch: 14]	cost = 0.00576675031
[Epoch: 15]	cost = 0.00654875953

Learning Finished!



MNIST CNN Code (Cont'd)

```
# Test model and check accuracy
with torch.no_grad():
    X_test = mnist_test.test_data.view(len(mnist_test), 1, 28, 28).float().to(device)
    Y_test = mnist_test.test_labels.to(device)

    prediction = model(X_test)
    correct_prediction = torch.argmax(prediction, 1) == Y_test
    accuracy = correct_prediction.float().mean()
    print('Accuracy:', accuracy.item())
```

```
Accuracy: 0.9869999885559082
```